



Improve Security of File System for Containers inBlackBox

Yongmao Luo, Zijian Zhang, Xincheng Xie





CONTENT

01

Background

02

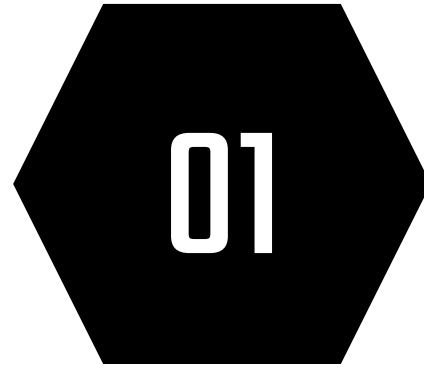
Proposal & Beta Impl

03

Demo

04

Result & Plan



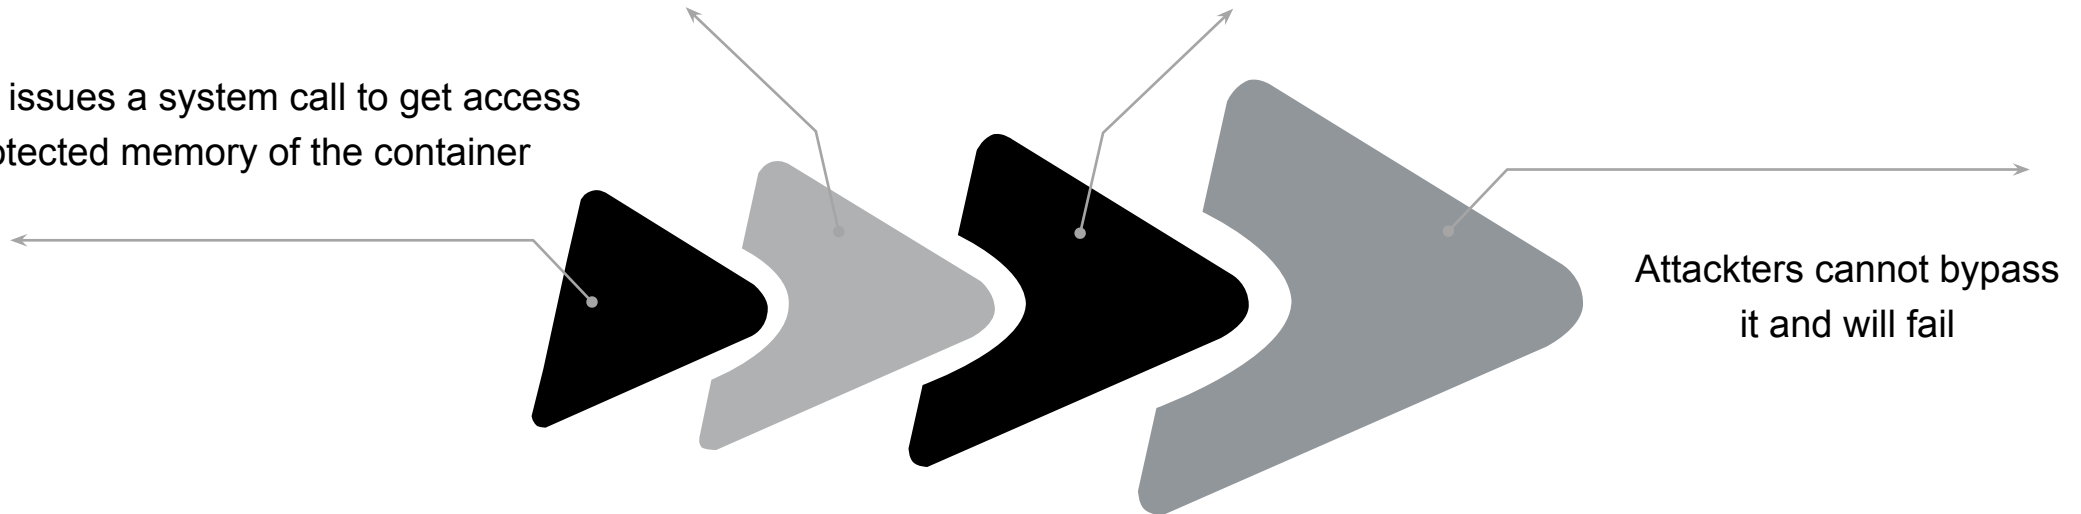
Background

Security provided by BlackBox

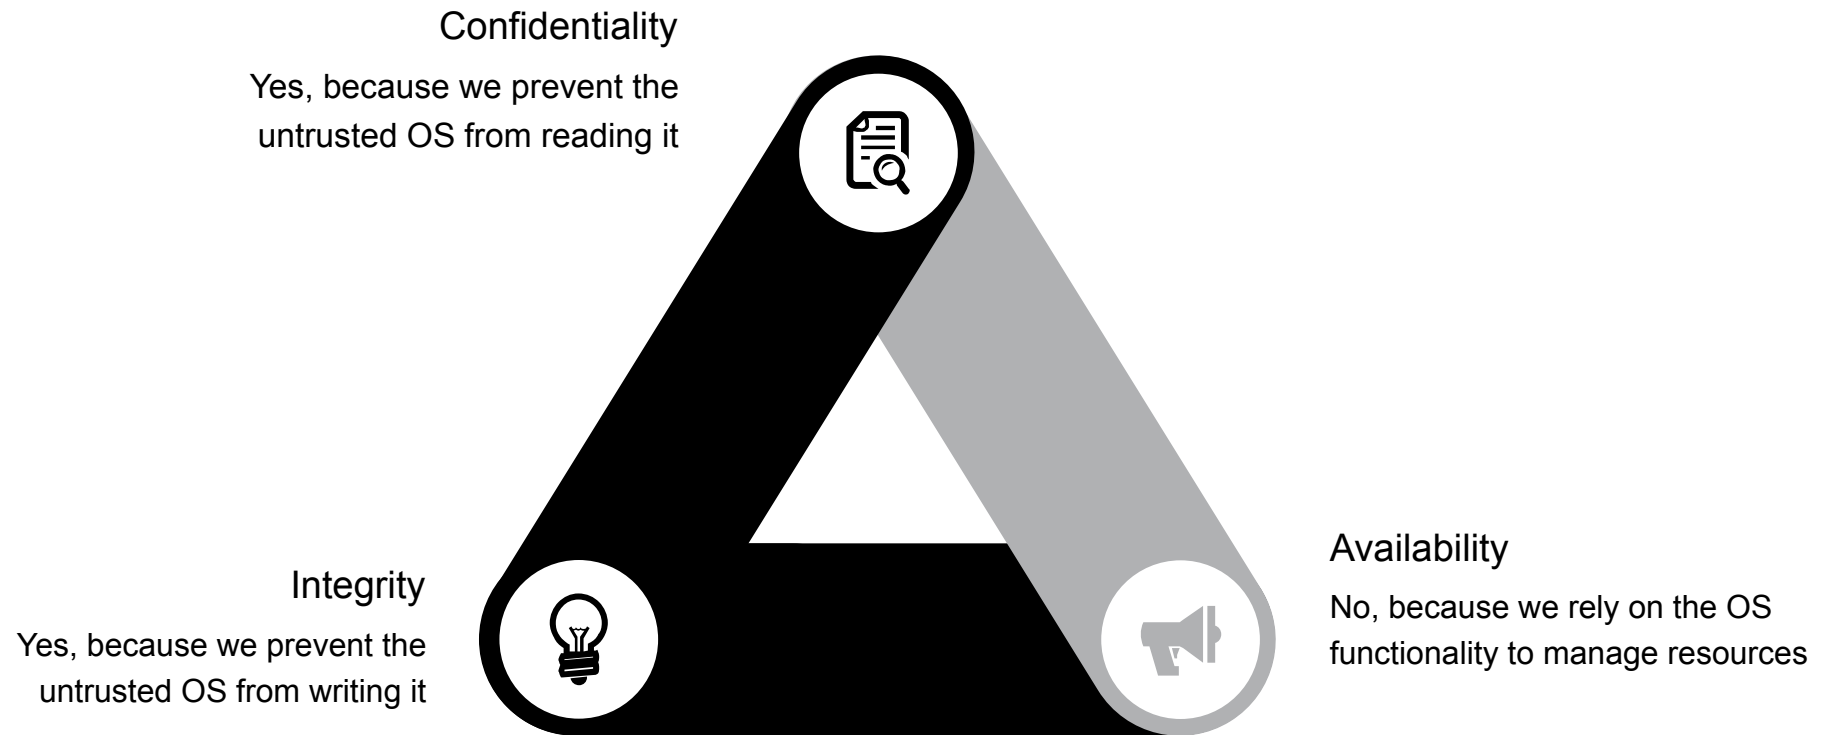
Attackers can refer to the protected memory directly through the NPT table of the OS

The EL2 translation by CSM enforce the admission control of access the memory

Attacker issues a system call to get access to protected memory of the container



Security provided by BlackBox





Files within a container can only be accessed through an OS's I/O facilities making access to a container's files inherently untrustworthy without additional protection. A userspace encrypted file system could potentially be used to provide transparent protection of file I/O, but this would likely signif-



02

Proposal & Beta Implementation



Encryption & Protected Memory

What to encrypt?

Newly created files

Newly created files by containers —
Private Files of Containers



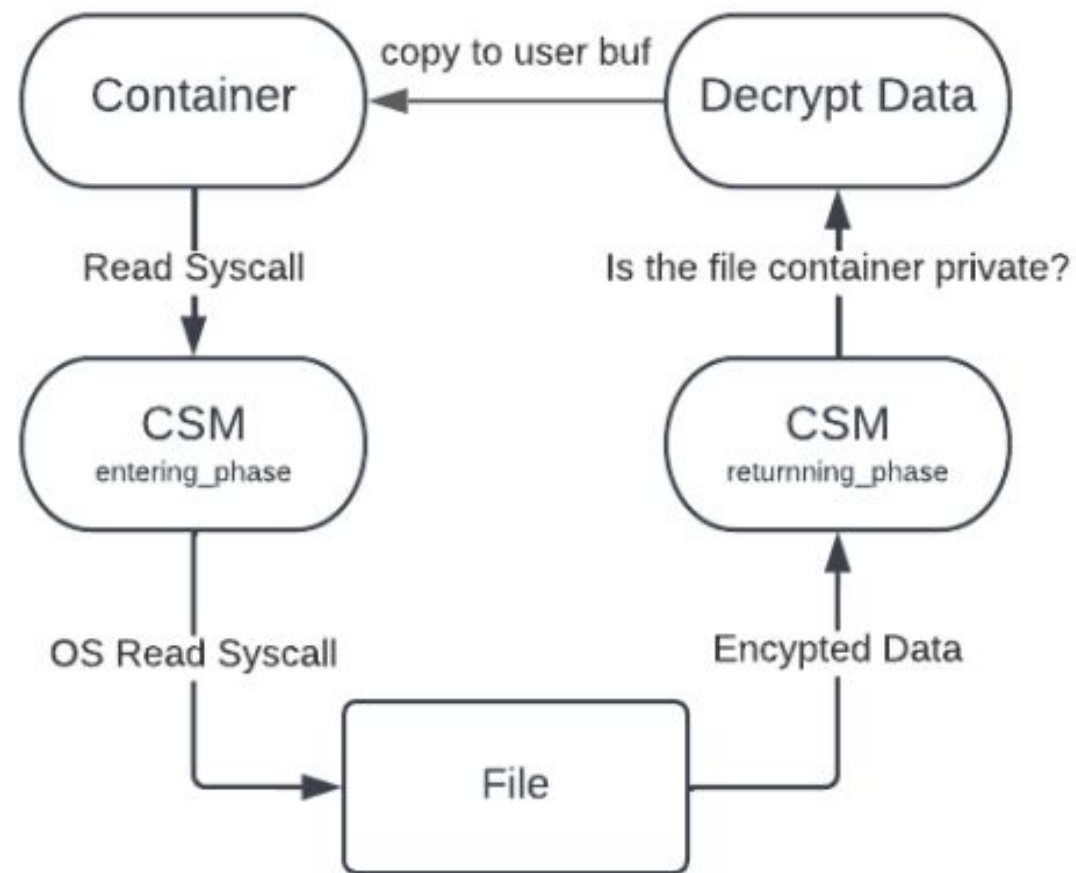


Figure 4: Read

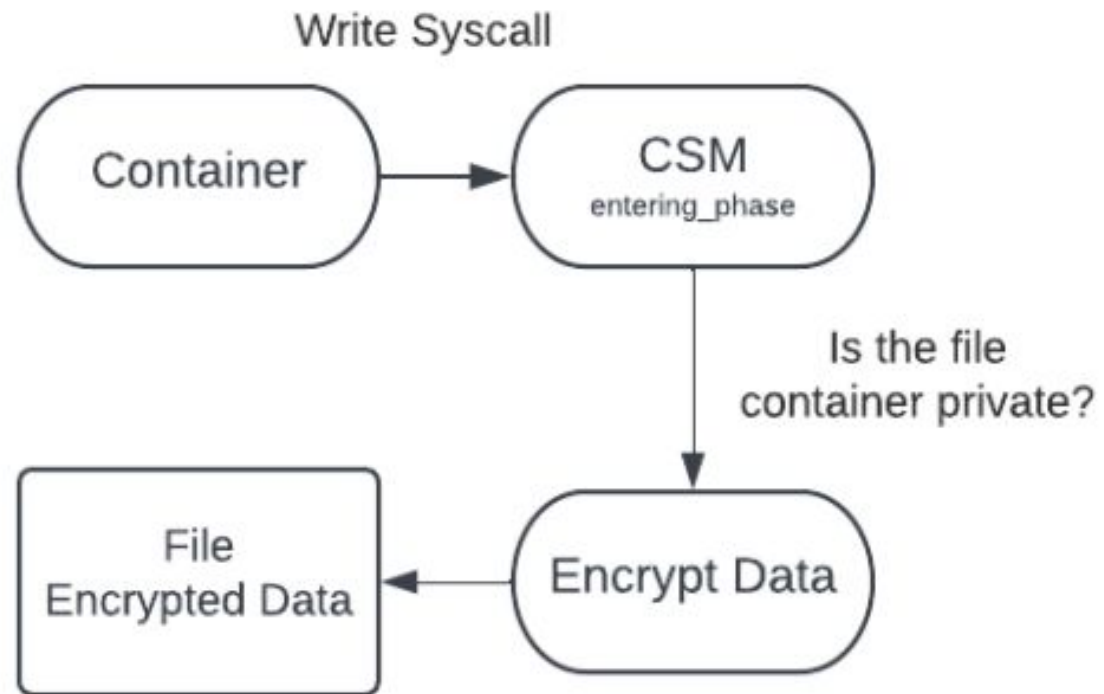
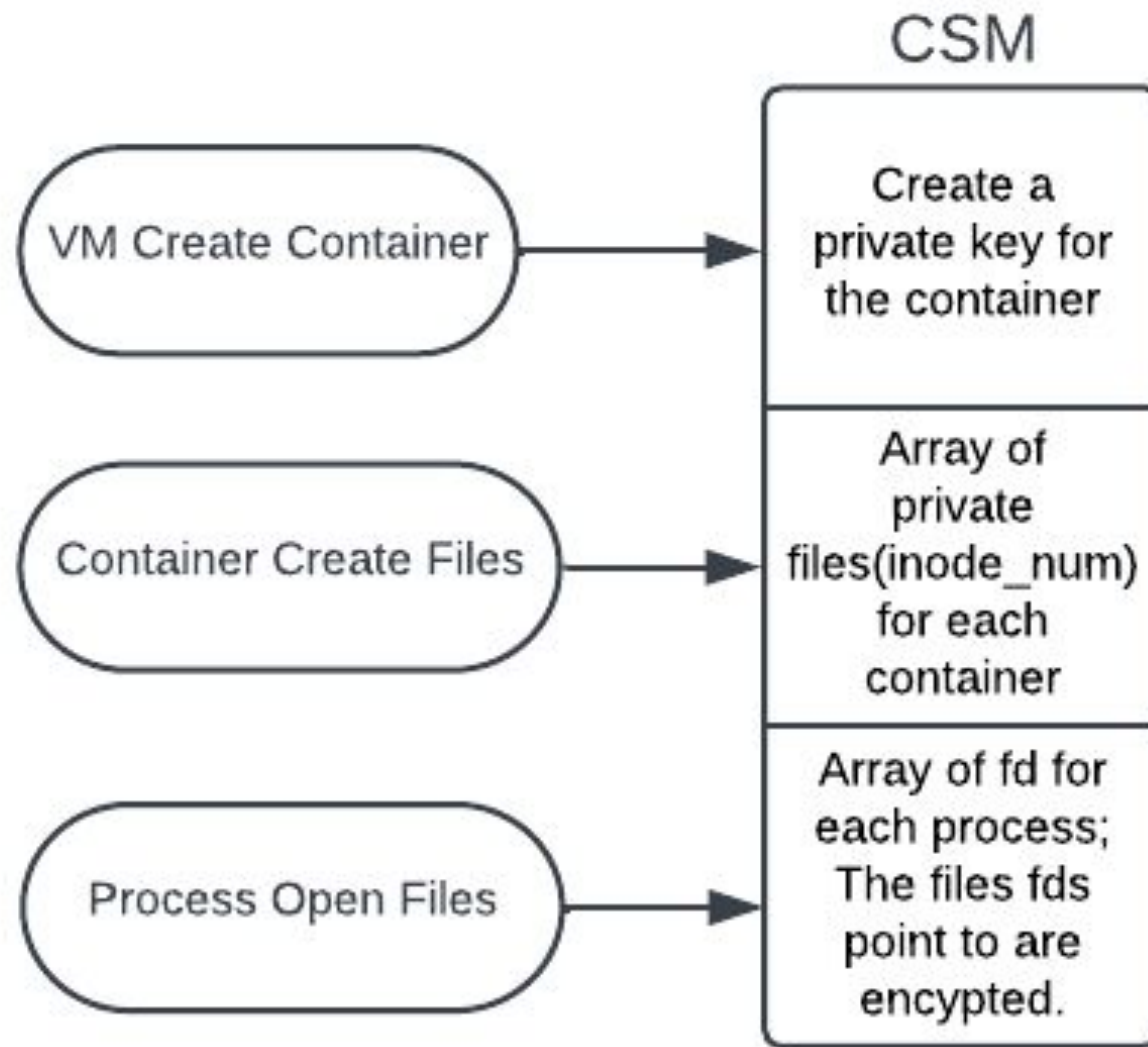



Figure 3: Write






```
int hvc_process_fork(uint32_t parent_eidx, uint32_t new_eidx, struct task_struct *child);
int hvc_process_exec(uint32_t caller_eidx, uint32_t new_eidx, struct task_struct *exec_tsk);
int hvc_process_clone(uint32_t caller_eidx, struct task_struct *thread);
void hvc_process_exit(void);
void hvc_restarting_syscall(int scno);

void hvc_settid(u64 tid_addr, u64 tid);

void hvc_file_create(uint64_t eidx, char *path, unsigned long i_ino);
void hvc_file_open(uint64_t eidx, char *path, unsigned long i_ino, int fd);
void hvc_file_unlink(uint64_t eidx, char *path, unsigned long i_ino);
void hvc_file_close(uint64_t eidx, char *path, int fd);
```

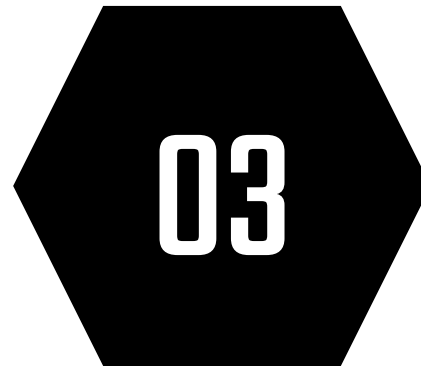


```
/* Negative dentry, just create the file */
if (!dentry->d_inode && (open_flag & O_CREAT)) {
    file->f_mode |= FMODE_CREATED;
    audit_inode_child(dir_inode, dentry, AUDIT_TYPE_CHILD_CREATE);
    if (!dir_inode->i_op->create) {
        error = -EACCES;
        goto out_dput;
    }
    error = dir_inode->i_op->create(dir_inode, dentry, mode,
                                open_flag & O_EXCL);
    if (error)
        goto out_dput;
    else if (eid_idx) {
        char *tmp_path = kmalloc(4096, GFP_KERNEL);
        char *path = dentry_path_raw(dentry, tmp_path, 4096);
        printk(KERN_DEBUG "lookup_open_create eid_idx = %lld file_path = %s i_ino = %ld\n",
               hvc_file_create(eid_idx, path, dentry->d_inode->i_ino);
        kfree(tmp_path);
    }
    fsnotify_create(dir_inode, dentry);
}
if (unlikely(create_error) && !dentry->d_inode) {
    error = create_error;
    goto out_dput;
}
```



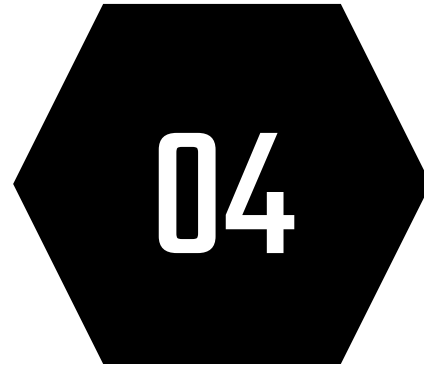

```
case HVC_FILE_CREATE: {
    u64 eidx = hr->regs[1];
    char *path = hr->regs[2];
    unsigned long i_ino = hr->regs[3];
    struct el2_enclave_task *e_task = get_current_enclave_task();
    struct el2_enclave_info* e_info = current_el2_enclave_info();
    int i;

    for (i = 0; i < MAX_ENCRYPTED_FILES; i++) {
        if (e_info->encrypted_ino_array[i] == -1) {
            e_info->encrypted_ino_array[i] = i_ino;
            printf("HVC_FILE_CREATE eidx = %lld file_path = %s i_ino = %ld\n", eidx, __el2_va(path), i_ino);
            printf("HVC_FILE_CREATE e_info->encrypted_ino_array[%d] = %d\n", i, i_ino);
            break;
        }
    }
    break;
}
```



Demo

<https://youtu.be/qwg2D7RWX-c>



Result & Plan

Result

1. Correctness.

- a. Processes within a container can correctly read the files
- b. Processes outside a container cannot correctly read the file content and fail to write.
- c. Compromised OS cannot correctly read the file content

2. Performance.

- a. file system benchmark: fio
- b. real application: LevelDB
- c. expectation: lower performance than unencrypted, but not significantly worse.



Basic Data

Structures:

1. track files belongings
2. storing ephemeral keys for each container

1

First Week

Hypercalls for manipulating such data structures

Improve memory

Utilization by deleting encrypted files generated by processes in container when closing the container

2

Second Week

Adding cryptography library to CSM to improve the security of encrypted files

Manipulate Read_write_wrapper to encrypt and decrypt data accordingly

3

Third Week

encrypt/decrypt in a more precise granularity

4

Fourth Week

Testing to ensure the functionality

Testing & final paper

5

Fifth Week





THANK YOU
For Your Listening